

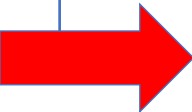
# UVA CS 6316: Machine Learning

## Lecture 14 Extra: More about Logistic Regression

Dr. Yanjun Qi

University of Virginia  
Department of Computer Science

# Today: Extra

- 
- ✓ Bayes Classifier and MAP Rule?
    - Bayes Classifier
    - Expected Prediction Error
    - 0-1 Loss function for Bayes Classifier

- ✓ Logistic regression
  - Parameter Estimation for LR

$$p(y|x) = \frac{e^{\beta x}}{1 + e^{\beta x}}$$

↓  
β

# Bayes classifiers

- Treat each feature attribute and the class label as random variables.
- Given a sample  $\mathbf{x}$  with attributes  $(x_1, x_2, \dots, x_p)$ :
  - Goal is to predict its class  $C$ .
  - Specifically, we want to find the value of  $C_i$  that maximizes  $p(C_i | x_1, x_2, \dots, x_p)$ .
- Can we estimate  $p(C_i | \mathbf{x}) = p(C_i | x_1, x_2, \dots, x_p)$  directly from data?

# Bayes classifiers

- Treat each feature attribute and the class label as random variables.
- Given a sample  $\mathbf{x}$  with attributes  $(x_1, x_2, \dots, x_p)$ :
  - Goal is to predict its class  $C$ .
  - Specifically, we want to find the value of  $C_i$  that maximizes  $p(C_i | x_1, x_2, \dots, x_p)$ .
- Can we estimate  $p(C_i | \mathbf{x}) = p(C_i | x_1, x_2, \dots, x_p)$  directly from data?

$\rightarrow \{C_1, C_2, \dots, C_L\}$

# Bayes classifiers

## → MAP classification rule

- Establishing a probabilistic model for classification

## → **MAP** classification rule

- **MAP**: **M**aximum **A** **P**osterior
- Assign  $x$  to  $c^*$  if

$$P(C = c^* | \mathbf{X} = \mathbf{x}) > P(C = c | \mathbf{X} = \mathbf{x})$$

$$\text{for } c \neq c^*, c = c_1, \dots, c_L$$

# Bayes classifiers

## → MAP classification rule

- Establishing a probabilistic model for classification

## → **MAP** classification rule

- **MAP**: **M**aximum **A** **P**osterior
- Assign  $x$  to  $c^*$  if

$$P(C = c^* | \mathbf{X} = \mathbf{x}) > P(C = c | \mathbf{X} = \mathbf{x})$$

$$c \neq c^*, c = c_1, \dots, c_L$$

$$\left\{ \begin{array}{l} P(C = c_1 | x) \\ P(C = c_2 | x) \\ P(C = c_3 | x) \end{array} \right\} \max \Rightarrow c_i$$

# Bayes Classifiers – MAP Rule

*Task:* Classify a new instance  $X$  based on a tuple of attribute values  $X = \langle X_1, X_2, \dots, X_p \rangle$  into one of the classes

$$c_{MAP} = \operatorname{argmax}_{c_j \in C} P(c_j | x_1, x_2, \dots, x_p)$$



WHY ?

MAP = Maximum A posteriori Probability

# 0-1 LOSS for Classification

if  $k = \ell$ ,  $L(k, \ell) = 0$


- Procedure for categorical output variable  $C$

if  $k \neq \ell$ ,  $L(k, \ell) = 1$

- Frequently, 0-1 loss function used:  $L(k, \ell)$

- $L(k, \ell)$  is the price paid for misclassifying an element from class  $C_k$  as belonging to class  $C_\ell$

→  $L \times L$  matrix



$C_1, C_2, \dots, C_L$



# Expected prediction error (EPE)

- Expected prediction error (EPE), with expectation taken w.r.t. the **joint distribution  $\Pr(C, X)$**

- $\Pr(C, X) = \Pr(C | X) \Pr(X)$

*→ e.g. 0-1 loss*

$$E_x(X)$$

$$E_x(g(X))$$

$$\text{EPE}(f) = E_{X,C}(L(C, f(X)))$$

$$= E_X \sum_{k=1}^L L[C_k, f(X)] \Pr(C_k | X)$$

Consider  
sample  
population  
distribution

$$EPE(f) = E_{\mathcal{X}, C} (L(C, f(\mathcal{X})))$$

$$E_C(C) = \sum_{i=1}^L C_i P(C_i)$$

$$= E_{\mathcal{X}} E_{C|\mathcal{X}} [L(C, f(\mathcal{X})) | \mathcal{X}]$$

Discrete RV's Expectation


$$= E_{\mathcal{X}} \sum_{k=1}^L L[C_k, f(\mathcal{X})] Pr(C_k | \mathcal{X})$$

$$\operatorname{argmin}_f EPE(f(\mathcal{X}))$$

$\Rightarrow$  Pointwise minimization when  $\mathcal{X} = x$

$$\Rightarrow \hat{f}(\mathcal{X} = x) = \operatorname{argmin}_{f(x) \in C} \sum_{k=1}^L L[C_k, f(x)] Pr(C_k | \mathcal{X} = x)$$

$$\Rightarrow \hat{f}(x) = \operatorname{argmax}_{C_k \in \left\{ \begin{matrix} C_1 \\ C_2 \\ C_3 \\ \vdots \\ C_L \end{matrix} \right\}} Pr(C_k | \mathcal{X} = x)$$



$$\begin{cases} p(C_1 | x) \\ p(C_2 | x) \\ \vdots \\ p(C_L | x) \end{cases}$$

# Expected prediction error (EPE)

$$\text{EPE}(f) = E_{X,C}(L(C, f(X))) = E_X \sum_{k=1}^K L(C_k, f(X)) \Pr(C_k | X)$$

Consider  
sample  
population  
distribution

- Pointwise minimization suffices

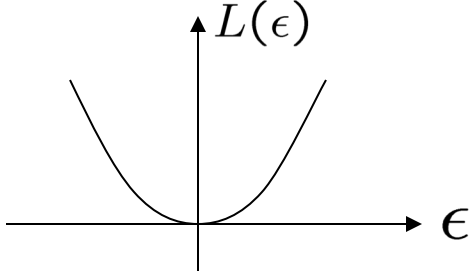
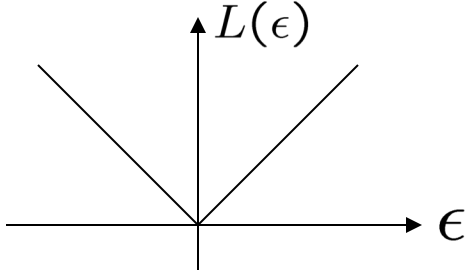
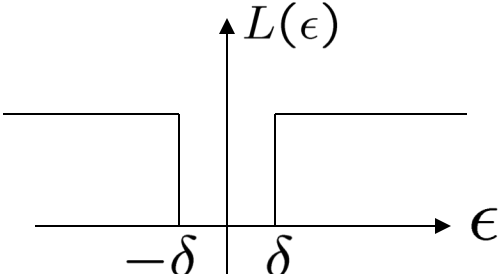
- $\rightarrow$  simply 
$$\hat{f}(X) = \operatorname{argmin}_{g \in \mathcal{C}} \sum_{k=1}^K L(C_k, g) \Pr(C_k | X = x)$$

**Bayes Classifier**

$$\hat{f}(X) = C_k \text{ if}$$

$$\Pr(C_k | X = x) = \max_{g \in \mathcal{C}} \Pr(g | X = x)$$

# SUMMARY: WHEN Expected prediction error (EPE) USES DIFFERENT LOSS

Loss Function	Estimator $\hat{f}(x)$
$L_2$ 	$EPE = E_{x,Y} (Y - f(x))^2$ $\hat{f}(x) = E[Y X = x]$
$L_1$ 	$\hat{f}(x) = \text{median}(Y X = x)$
$0-1$ 	$\hat{f}(x) = \arg \max_Y P(Y X = x)$ <p>(Bayes classifier / MAP)</p>

# Today: Extra

- ✓ Why Bayes Classification – MAP Rule?
  - Expected Prediction Error
  - 0-1 Loss function for Bayes Classifier

- ✓ Logistic regression

- 
- Parameter Estimation for LR

$$p(y|x) = \frac{e^{\beta x}}{1 + e^{\beta x}}$$

$\Downarrow$   
 $\beta$

# Newton's method for optimization

- The most basic **second-order** optimization algorithm
- Updating parameter with

$$\text{GD: } \theta_{k+1} = \theta_k - \alpha g_k$$

$$\text{Newton: } \theta_{k+1} = \theta_k - \underbrace{\mathbf{H}_K^{-1}}_{p \times p} \mathbf{g}_k$$

$p \times 1$

# Review: Hessian Matrix / $n=2$ case

Singlevariate  $\rightarrow$  multivariate

- 1<sup>st</sup> derivative to gradient,
- 2<sup>nd</sup> derivative to Hessian

$$f(x, y)$$

$$g = \nabla f = \begin{pmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{pmatrix}$$

$$H = \begin{pmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} \end{pmatrix}$$

# Review: Hessian Matrix

Suppose that  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is a function that takes a vector in  $\mathbb{R}^n$  and returns a real number. Then the **Hessian** matrix with respect to  $x$ , written  $\nabla_x^2 f(x)$  or simply as  $H$  is the  $n \times n$  matrix of partial derivatives,

$$\nabla_x^2 f(x) \in \mathbb{R}^{n \times n} = \begin{bmatrix} \frac{\partial^2 f(x)}{\partial x_1^2} & \frac{\partial^2 f(x)}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f(x)}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f(x)}{\partial x_2 \partial x_1} & \frac{\partial^2 f(x)}{\partial x_2^2} & \cdots & \frac{\partial^2 f(x)}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f(x)}{\partial x_n \partial x_1} & \frac{\partial^2 f(x)}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f(x)}{\partial x_n^2} \end{bmatrix}.$$



# Newton's method for optimization

- Making a quadratic/second-order Taylor series approximation

$$\hat{f}_{quad}(\boldsymbol{\theta}) = f(\boldsymbol{\theta}_k) + \mathbf{g}_k^T (\boldsymbol{\theta} - \boldsymbol{\theta}_k) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_k)^T \mathbf{H}_k (\boldsymbol{\theta} - \boldsymbol{\theta}_k)$$

Finding the minimum solution of the above right quadratic approximation (quadratic function minimization is easy !)

$$\hat{f}(\theta) = f(\theta_k) + g_k^T (\theta - \theta_k) + \underbrace{\frac{1}{2} (\theta - \theta_k)^T H_k (\theta - \theta_k)}_{\Downarrow}$$

$$\frac{1}{2} (\theta^T H_k \theta - 2 \theta^T H_k \theta_k + \theta_k^T H_k \theta_k)$$


$$\frac{\partial \hat{f}(\theta)}{\partial \theta} = 0 + g_k + \underbrace{\frac{2}{2} H_k \theta - \frac{2}{2} H_k \theta_k}_{\text{see p24 handout}} = 0$$

$$g_k + H_k (\theta - \theta_k) = 0$$


$$\Rightarrow \theta = \theta_k - H_k^{-1} g_k$$

where  $H_k \in \mathbb{R}^{p \times p}$   
 $g_k \in \mathbb{R}^p$


# Newton's Method / second-order Taylor series approximation




# Newton's Method / second-order Taylor series approximation



# Newton's Method / second-order Taylor series approximation



# Newton's Method / second-order Taylor series approximation



# Newton's Method

- At each step:

$$\theta_{k+1} = \theta_k - \frac{f'(\theta_k)}{f''(\theta_k)}$$

$$\theta_{k+1} = \theta_k - H^{-1}(\theta_k) \nabla f(\theta_k)$$

- Requires 1<sup>st</sup> and 2<sup>nd</sup> derivatives
- Quadratic convergence
- ➔ However, finding the inverse of the Hessian matrix is often expensive

# Newton vs. GD for optimization

- **Newton:** a quadratic/second-order Taylor series approximation

$$\hat{f}_{quad}(\boldsymbol{\theta}) = f(\boldsymbol{\theta}_k) + \mathbf{g}_k^T (\boldsymbol{\theta} - \boldsymbol{\theta}_k) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_k)^T \mathbf{H}_k (\boldsymbol{\theta} - \boldsymbol{\theta}_k)$$

$\Rightarrow \boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \frac{1}{H(\boldsymbol{\theta}_k)} g(\boldsymbol{\theta}_k)$

Finding the minimum solution of the above right quadratic approximation (quadratic function minimization is easy !)

$$\hat{f}_{quad}(\boldsymbol{\theta}) = f(\boldsymbol{\theta}_k) + \mathbf{g}_k^T (\boldsymbol{\theta} - \boldsymbol{\theta}_k) + \frac{1}{2} (\boldsymbol{\theta} - \boldsymbol{\theta}_k)^T \frac{1}{\alpha} (\boldsymbol{\theta} - \boldsymbol{\theta}_k)$$

$\Downarrow \boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \alpha g(\boldsymbol{\theta}_k)$




# Comparison

- Newton's method vs. Gradient descent

A comparison of gradient descent (green) and Newton's method (red) for minimizing a function (with small step sizes).

Newton's method uses curvature information to get a more direct route ...



# MLE for Logistic Regression Training

Let's fit the logistic regression model for  $K=2$ , i.e., number of classes is 2

Training set:  $(x_i, y_i), i=1, \dots, N$

For Bernoulli distribution

$$p(y | x)^y (1 - p)^{1-y}$$

(conditional)  
Log-likelihood:

How?

$$\begin{aligned} l(\beta) &= \sum_{i=1}^N \{\log \Pr(Y = y_i | X = x_i)\} \\ &= \sum_{i=1}^N y_i \log(\Pr(Y = 1 | X = x_i)) + (1 - y_i) \log(\Pr(Y = 0 | X = x_i)) \\ &= \sum_{i=1}^N \left( y_i \log \frac{\exp(\beta^T x_i)}{1 + \exp(\beta^T x_i)} + (1 - y_i) \log \frac{1}{1 + \exp(\beta^T x_i)} \right) \\ &= \sum_{i=1}^N (y_i \beta^T x_i - \log(1 + \exp(\beta^T x_i))) \end{aligned}$$

$x_i$  are  $(p+1)$ -dimensional input vector with leading entry 1

$\beta$  is a  $(p+1)$ -dimensional vector

10/23/19 Dr. Pradyumn Kumar MCS We want to **maximize** the log-likelihood in order to estimate  $\beta$

$$l(\beta) = \sum_{i=1}^N \{\log \Pr(Y = y_i | X = x_i)\}$$

$y_i$

$p(y_i=1|x)$

$$\begin{aligned} & \log \{ \Pr(Y = y_i | X = x_i) = p(y_i | x_i) \} \Rightarrow \begin{matrix} y_i = 1 \\ y_i = 0 \end{matrix} \\ & = \log \{ p(y_i=1|x)^{y_i} (1 - p(y_i=1|x))^{1-y_i} \} \\ & = y_i \log p(y_i=1|x) + (1-y_i) \log (1 - p(y_i=1|x)) \end{aligned}$$

# Newton-Raphson for LR (optional)

$$\frac{\partial l(\beta)}{\partial \beta} = \sum_{i=1}^N (y_i - \frac{\exp(\beta^T x)}{1 + \exp(\beta^T x)}) x_i = 0$$

(p+1) Non-linear equations to solve for (p+1) unknowns

Vector  $\beta$

Solve by Newton-Raphson method:

$$\beta^{new} \leftarrow \beta^{old} - [(\frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T})]^{-1} \frac{\partial l(\beta)}{\partial \beta},$$

$$\text{where, } (\frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T}) = - \sum_{i=1}^N x_i x_i^T (\frac{\exp(\beta^T x_i)}{1 + \exp(\beta^T x_i)}) (\frac{1}{1 + \exp(\beta^T x_i)})$$

minimizes a quadratic approximation to the function we are really interested in.

$$\theta_{k+1} = \theta_k - \mathbf{H}_K^{-1} \mathbf{g}_k$$

$p(x_i; \beta)$

$1 - p(x_i; \beta)$

# Newton-Raphson for LR...

$$\frac{\partial l(\beta)}{\partial \beta} = \sum_{i=1}^N \left( y_i - \frac{\exp(\beta^T x)}{1 + \exp(\beta^T x)} \right) x_i = X^T (y - p)$$

$\rightarrow p(y=1|x) = \frac{e^{\beta^T x}}{1 + e^{\beta^T x}}$

$$\left( \frac{\partial^2 l(\beta)}{\partial \beta \partial \beta^T} \right) = -X^T W X$$

So, NR rule becomes:

$$\beta^{new} \leftarrow \beta^{old} + (X^T W X)^{-1} X^T (y - p),$$

$$X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_N^T \end{bmatrix}_{N \times (p+1)}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}_{N \times 1}, \quad p = \begin{bmatrix} \exp(\beta^T x_1) / (1 + \exp(\beta^T x_1)) \\ \exp(\beta^T x_2) / (1 + \exp(\beta^T x_2)) \\ \vdots \\ \exp(\beta^T x_N) / (1 + \exp(\beta^T x_N)) \end{bmatrix}_{N \times 1}$$

$X : N \times (p+1)$  matrix of  $x_i$

$y : N \times 1$  matrix of  $y_i$

$p : N \times 1$  matrix of  $p(x_i; \beta^{old})$

$W : N \times N$  diagonal matrix of  $p(x_i; \beta^{old})(1 - p(x_i; \beta^{old}))$

$$\left( \frac{\exp(\beta^T x_i)}{(1 + \exp(\beta^T x_i))} \right) \left( 1 - \frac{1}{(1 + \exp(\beta^T x_i))} \right)$$

# Newton-Raphson for LR...

- Newton-Raphson

$$\begin{aligned} - \beta^{new} &= \beta^{old} + (X^T W X)^{-1} X^T (y - p) \\ &= (X^T W X)^{-1} X^T W (X \beta^{old} + W^{-1} (y - p)) \\ &= (X^T W X)^{-1} X^T W z \end{aligned}$$

Re expressing  
Newton step as  
weighted least  
square step

- Adjusted response


$$z = X \beta^{old} + W^{-1} (y - p)$$

- Iteratively reweighted least squares (IRLS)

$$\beta^{new} \leftarrow \arg \min_{\beta} (z - X \beta^T)^T W (z - X \beta^T)$$

$$\leftarrow \arg \min_{\beta} (y - p)^T W^{-1} (y - p)$$

# Binary → Multinoulli Logistic Regression Model



$$p(G|x)$$

Directly models the posterior probabilities as the output of regression

$$\Pr(G = k | X = x) = \frac{\exp(\beta_{k0} + \beta_k^T x)}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)}, \quad k = 1, \dots, K-1$$

$$\Pr(G = K | X = x) = \frac{1}{1 + \sum_{l=1}^{K-1} \exp(\beta_{l0} + \beta_l^T x)}$$

$x$  is  $p$ -dimensional input vector

$\beta_k^T$  is a  $p$ -dimensional vector for each class  $k$


Total number of parameters is  $(K-1)(p+1)$

$\beta_{k0}, \vec{\beta}_k, k=1, 2, \dots, K-1$


Note that the class boundaries are linear

# Binary $\rightarrow$ Multinoulli Logistic Regression Model

(e.g.  $k=6$ )



$$p(y=1|x) = \frac{e^{\beta_1 x}}{1 + e^{\beta_1 x}}$$
$$p(y=0|x) = \frac{1}{1 + e^{\beta_1 x}}$$



$$\frac{e^{\beta_k^T x}}{1 + e^{\beta_1^T x} + e^{\beta_2^T x} + \dots}$$

e.g.

$$\ln \frac{P(G=k|x)}{P(G=k'|x)} = 0 \Rightarrow \text{linear}$$
$$\beta_{k0} + \beta_k^T x$$

Note that the class boundaries are **linear**



# References

- ❑ Prof. Tan, Steinbach, Kumar's "Introduction to Data Mining" slide
- ❑ Prof. Andrew Moore's slides
- ❑ Prof. Eric Xing's slides
- ❑ Prof. Ke Chen NB slides
- ❑ Hastie, Trevor, et al. *The elements of statistical learning*. Vol. 2. No. 1. New York: Springer, 2009.