

# UVA CS 6316: Machine Learning

## Lecture 18c: More and Extra about Boosting

Dr. Yanjun Qi

University of Virginia  
Department of Computer Science

## Boosting:

- Learners are ordered: Each learner tries to reduce error (residual) on “hard” examples (those misclassified by earlier learners).
- ADABOOST: weight hard samples more;
- GRADIENT BOOST: use residual to train later models. Reduces bias and possibly variance compared to base learners.
- Gradient-boosted decision trees (GBDT) often gives state-of-the-art performance on simple classification tasks, e.g. XGBOOST.
- Neural networks are used fairly often with bagging, but rarely with boosting.
- Decision trees work well in both bagging and boosting.

# Boosting

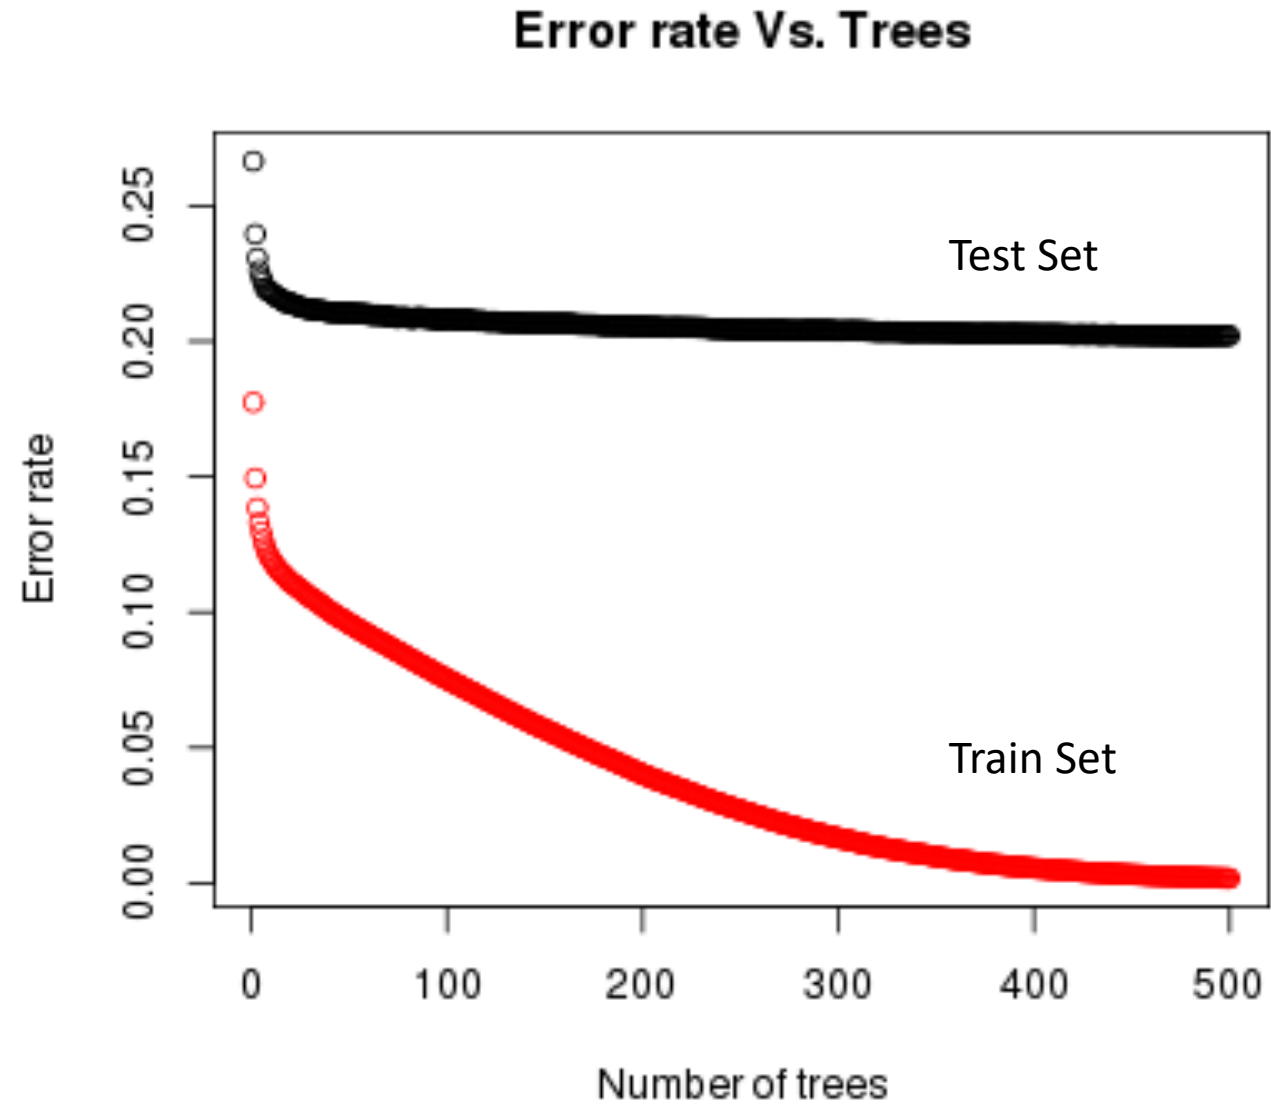
- Sequential algorithm where at each step, a weak learner is trained based on the results of the previous learner.
- Two main types:
  - **Adaptive Boosting:** Reweight datapoints based on performance of last weak learner. Focuses on points where previous learner had trouble. Example: AdaBoost.
  - **Gradient Boosting:** Train new learner on residuals of overall model. Constitutes gradient boosting because approximating the residual and adding to the previous result is essentially a form of gradient descent. Example: XGBoost.

# Gradient Boosting

# XGBoost

- XGBoost is a very efficient Gradient Boosting Decision Tree implementation with some interesting features:
  - **Regularization:** Can use L1 or L2 regularization.
  - **Handling sparse data:** Incorporates a sparsity-aware split finding algorithm to handle different types of sparsity patterns in the data.
  - **Weighted quantile sketch:** Uses distributed weighted quantile sketch algorithm to effectively handle weighted data.
  - **Block structure for parallel learning:** Makes use of multiple cores on the CPU, possible because of a block structure in its system design. Block structure enables the data layout to be reused.
  - **Cache awareness:** Allocates internal buffers in each thread, where the gradient statistics can be stored.
  - **Out-of-core computing:** Optimizes the available disk space and maximizes its usage when handling huge datasets that do not fit into memory.

# XGBoost (an example performance figure )



# Gradient Boosting

- Task is to estimate target continuous function  $F(x)$ . We measure goodness of estimation with loss function  $L(y, F(x))$ .
- Gradient boosting assumes that:
- $F(x) = \alpha_0 + \alpha_1 h_1(x) + \dots + \alpha_M h_M(x)$
- Basic Gradient boosting workflow:

1. Initialize  $F_0(x) = \alpha_0$
2. Estimate  $\alpha_m$  and  $h_m(x)$  such that:

$$L(y, F_{m-1}(x) + \alpha_m h_m(x)) < L(y, F_{(m-1)}(x))$$

3. Update  $F_m(x) = F_{m-1}(x) + \alpha_m h_m(x)$
4. Repeat from 2, M times.

# Gradient Boosting

$$L(y, F_{m-1}(x) + \alpha_m h_m(x)) < L(y, F_{(m-1)}(x))$$



If we can find a vector  $r_m$  that we can plug in here to make this equation true, we can train a basic learner  $h_m(x)$  to predict  $r_m$  from  $x$ !

We are basically searching for a vector that points to the direction that reduces our loss... does that sound familiar?

**Gradient descent!**



# Gradient Boosting

- By solving a simple 1D optimization problem, we could also find the optimal  $\alpha_m$  for each step, by computing:
- $\alpha_m = \operatorname{argmin}_{\gamma} L(y, F_{m-1}(x) + \gamma h_m(x))$
- This gives us an updated Gradient Boosting algorithm:
  1. Initialize  $F_0(x) = \alpha_0$
  2. Compute negative gradient per observation:  $r_{m_i} = -\frac{\partial L(y_i, F_{m-1}(x_i))}{\partial F_{m-1}(x_i)}$
  3. Train base learner  $h_m(x)$  on the gradients
  4. Compute  $\alpha_m$  with line search strategy
  5. Update  $F_m(x) = F_{m-1}(x) + \alpha_m h_m(x)$
  6. Repeat from 2, M times.

# Gradient Boosting

- Where do the residuals come in?
- If we consider Mean Squared Error as our loss function, the per-observation gradient is:

- $$\frac{\partial L(y_i, F_m(x_i))}{\partial F_m(x_i)} = \frac{\partial \left( \frac{1}{2n} \sum_i (y_i - F_m(x_i))^2 \right)}{\partial F_m(x_i)} = \frac{\partial \left( \frac{1}{2} (y_i - F_m(x_i))^2 \right)}{\partial F_m(x_i)} = y_i - F_m(x_i)$$

- The derivation we found before works with any loss function.

# Gradient Tree Boosting

- When dealing with decision trees, we can take the concept further by selecting a specific  $\alpha_m$  for each of the tree's regions. The output of a tree is:
- $h_m(x) = \sum_{j=1}^{J_m} b_{jm} \mathbf{1}_{R_{jm}}(x)$
- The model update rule becomes:
- $F_m(x) = F_{m-1}(x) + \sum_{j=1}^{J_m} \alpha_{jm} \mathbf{1}_{R_{jm}}(x)$
- $\alpha_{jm} = \operatorname{argmin}_{\gamma} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma)$

$J_m$ : Disjoint regions  
partitioned by the tree

$R_{jm}$ : Number  
of leaves

# XGBoost

- Three main forms of gradient boosting are supported:
- **Gradient Boosting** algorithm, as we defined above.
- **Stochastic Gradient Boosting** with sub-sampling at the row, column and column per split levels.
  - Random procedure where we subsample observations and features
- **Regularized Gradient Boosting** with both L1 and L2 regularization.
  - add a regularization term to the loss function that we are optimizing:

$$L_R(y, F(x)) = L(y, F(x)) + \Omega(F)$$

$$\text{Where } \Omega(F) = \gamma T + \frac{1}{2} \lambda \|w\|^2$$

T: Number of leaves

W: Leaf weights: prediction of each leaf

credit: Camilo Fosco

# XGBoost

- Remember, we still want to find the tree structure that minimizes our loss, which means best score structure. Doing this for all possible tree structures is unfeasible.
- A greedy algorithm that starts from a single leaf and iteratively adds branches to the tree is used instead.

# XGBoost

- XGBoost adds multiple other important advancements that make it state of the art in several industrial applications.
- In practice:
  - Can take a while to run if you don't set the `n_jobs` parameter correctly
  - Defining the `eta` parameter (analogous to learning rate) and `max_depth` is crucial to obtain good performance.
  - Alpha parameter controls L1 regularization, can be increased on high dimensionality problems to increase run time.

# XGBoost

- General approach to parameter tuning:
  - Cross-validate **learning rate**.
  - Determine the **optimum number of trees for this learning rate**. XGBoost can perform cross-validation at each boosting iteration for this, with the “cv” function.
  - **Tune tree-specific parameters** (max\_depth, min\_child\_weight, gamma, subsample, colsample\_bytree) for chosen learning rate and number of trees.
  - Tune **regularization parameters** (lambda, alpha).

# LGBM

- Stands for Light Gradient Boosted Machines. It is a library for training GBMs developed by Microsoft, and it competes with XGBoost.
- Extremely **efficient implementation**.
- Usually much faster than XGBoost with low hit on accuracy.
- Main contributions are two novel techniques to speed up split analysis: **Gradient based one-side sampling** and **Exclusive Feature Building**.
- Leaf-wise tree growth vs level-wise tree growth of XGBoost.

